


```
<configuration>
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://10.0.195.85:9000</value>
</property>
</configuration>
hadoop@aminpri-2-of-3:~/hadoop$
```

```
<configuration>
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://10.0.195.85:9000</value>
</property>
</configuration>
hadoop@aminpri-3-of-3:~/hadoop$
```

Step5: Edit the hdfs-site.xml file. Change the NameNode and DataNode directory paths.

Namenode - master

Datanode - slave nodes

```
<configuration>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:///home/hadoop/hadoopdata/hdfs/namenode</value>
  </property>

  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:///home/hadoop/hadoopdata/hdfs/datanode</value>
  </property>
</configuration>
hadoop@aminpri-1-of-3:~$
```

```
<configuration>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>file:///home/hadoop/hadoopdata/hdfs/datanode</value>
</property>
</configuration>
hadoop@aminpri-2-of-3:~/hadoop$
```

```
<configuration>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>file:///home/hadoop/hadoopdata/hdfs/datanode</value>
</property>
</configuration>
hadoop@aminpri-3-of-3:~/hadoop$
```

Step6: Edit the mapred-site.xml file. This would be the same for all nodes.

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
  <property>
    <name>mapreduce.admin.user.env</name>
    <value>HADOOP_MAPRED_HOME=$HADOOP_COMMON_HOME</value>
  </property>
  <property>
    <name>yarn.app.mapreduce.am.env</name>
    <value>HADOOP_MAPRED_HOME=$HADOOP_COMMON_HOME</value>
  </property>
</configuration>
hadoop@aminpri-1-of-3:~$
```

Step7: Edit the yarn-site.xml file. Add the ip of the master node as resource manager for all nodes.

```
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
  <description>The hostname of the RM.</description>
  <name>yarn.resourcemanager.hostname</name>
  <value>10.0.195.85</value>
</property>
<property>
  <description>The address of the applications manager interface in the RM.</description>
  <name>yarn.resourcemanager.address</name>
  <value>10.0.195.85:8032</value>
</property>
</configuration>
hadoop@aminpri-1-of-3:~$
```

```
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
  <description>The hostname of the RM.</description>
  <name>yarn.resourcemanager.hostname</name>
  <value>10.0.195.85</value>
</property>
<property>
  <description>The address of the applications manager interface in the RM.</description>
  <name>yarn.resourcemanager.address</name>
  <value>10.0.195.85:8032</value>
</property>
</configuration>
hadoop@aminpri-2-of-3:~/hadoop$
```

```
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
  <description>The hostname of the RM.</description>
  <name>yarn.resourcemanager.hostname</name>
  <value>10.0.195.85</value>
</property>
<property>
  <description>The address of the applications manager interface in the RM.</description>
  <name>yarn.resourcemanager.address</name>
  <value>10.0.195.85:8032</value>
</property>
</configuration>
hadoop@aminpri-3-of-3:~/hadoop$
```

Step8: Add the ip address of the slave nodes in the workers file of the master node. No change for slave nodes

```
hadoop@aminpri-1-of-3:~/hadoop/etc/hadoop$ cat workers
10.0.195.85
10.0.195.92
10.0.195.197
hadoop@aminpri-1-of-3:~/hadoop/etc/hadoop$
```

```
hadoop@aminpri-2-of-3:~/hadoop/etc/hadoop$ cat workers
localhost
hadoop@aminpri-2-of-3:~/hadoop/etc/hadoop$ |
```

```
hadoop@aminpri-3-of-3:~/hadoop/etc/hadoop$ cat workers
localhost
hadoop@aminpri-3-of-3:~/hadoop/etc/hadoop$
```

Step9: Format the Namenode as a hadoop user (only on Master) and run the command start-all.sh (only on Master).

Check the status of all nodes using jps.

You should see the NameNode, SecondaryNameNode, and ResourceManager daemons running on your master, and DataNode and NodeManager daemons on your slaves.

```
hadoop@aminpri-1-of-3:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [aminpri-1-of-3.js2local]
Starting datanodes
Starting secondary namenodes [aminpri-1-of-3]
Starting resourcemanager
Starting nodemanagers
hadoop@aminpri-1-of-3:~$ jps
11329 Jps
10531 DataNode
10886 ResourceManager
10392 NameNode
11000 NodeManager
10698 SecondaryNameNode
hadoop@aminpri-1-of-3:~$
```

```
hadoop@aminpri-2-of-3:~$ jps
132131 DataNode
132228 NodeManager
132349 Jps
hadoop@aminpri-2-of-3:~$
```

```
hadoop@aminpri-3-of-3:~$ jps
132484 Jps
132266 DataNode
132363 NodeManager
hadoop@aminpri-3-of-3:~$
```

A 3-node cluster has been setup successfully.

References:

1. *Installation of Hadoop on Ubuntu 2023-*
<https://sanjayshankar1428.medium.com/in-the-world-of-big-data-managing-and-processing-vast-amounts-of-information-efficiently-is-2ffdc0204357>
2. *Hadoop Multi Node Cluster Setup* - https://www.youtube.com/watch?v=_iP2Em-5Abw

NY Parking Violations

1. When are tickets most likely to be issued?

```
+-----+-----+
|violation_time|count|
+-----+-----+
|      0836A|20112|
|      0839A|19531|
|      0840A|19436|
|      0838A|19433|
|      0906A|19412|
|      1139A|19255|
|      1140A|19154|
|      1141A|19135|
|      0841A|18987|
|      1142A|18936|
|      1145A|18917|
|      0837A|18893|
|      0842A|18806|
|      1143A|18738|
|      0910A|18698|
|      1138A|18679|
|      0845A|18648|
|      0908A|18648|
|      0909A|18604|
|      1136A|18541|
|      1144A|18524|
|      0907A|18464|
|      0843A|18363|
|      1146A|18324|
|      1147A|18313|
|      0844A|18185|
|      1137A|17977|
|      0846A|17964|
|      1148A|17963|
|      0912A|17929|
+-----+-----+
only showing top 30 rows
```

2. What are the most common years and types of cars to be ticketed?

```
+-----+-----+-----+
|vehicle_type|vehicle_year|violation_count|
+-----+-----+-----+
|      SUBN|      2021|      468828|
|      SUBN|      2022|      452377|
|      SUBN|      2023|      447136|
|      SUBN|      2019|      345021|
|      SUBN|      2020|      343283|
|      SUBN|      2018|      275702|
|      SUBN|      2017|      226828|
|      SUBN|      2016|      186233|
|      SUBN|      2015|      180925|
|      4DSD|      2017|      155318|
|      4DSD|      2019|      152339|
|      4DSD|      2018|      146010|
|      SUBN|      2014|      142950|
|      4DSD|      2020|      138040|
|      4DSD|      2023|      134263|
|      4DSD|      2021|      134002|
|      4DSD|      2022|      131561|
|      SUBN|      2013|      130515|
|      4DSD|      2015|      129270|
|      4DSD|      2016|      126511|
|      SUBN|      2011|      117729|
|      SUBN|      2012|      117674|
|      4DSD|      2013|      108203|
|      4DSD|      2014|      102357|
|      SUBN|      2008|       93048|
|      SUBN|      2010|       88975|
|      4DSD|      2012|       87130|
|      SUBN|      2007|       77786|
|      4DSD|      2010|       74395|
|      VAN|      2019|       73346|
+-----+-----+-----+
only showing top 30 rows
```

3. Where are tickets most commonly issued?

Number_of_tickets	Violation_Location
276203	19
213205	114
207636	6
189589	13
178348	14
153765	109
148286	1
147809	18
142074	9
135832	115
116439	61
115903	66
115747	20
109812	112
107721	70
104404	84
104246	103
103097	52
102733	108
98620	46
94933	10
88065	110
84947	90
84910	102
83649	43
83088	104
81657	24
80699	62
76885	44
76378	17

only showing top 30 rows

4. Which color of the vehicle is most likely to get a ticket?

Vehicle_Color	count
GY	2086349
WH	1924604
BK	1821703
NULL	1015118
BL	688918
WHITE	610935
BLACK	401993
RD	393388
GREY	303176
BLUE	140721
GR	134699
SILVE	134375
BROWN	129885
RED	116436
BLK	89406
TN	70852
BR	68615
YW	65868
GRY	64505
WHI	54907
GL	39134
OTHER	33184
MR	30204
GREEN	26247
OR	24657
SIL	22166
GRAY	21722
BLU	20578
YELLOW	15621
LTGY	12880

only showing top 30 rows

5. Given a Black vehicle parking illegally at 34510, 10030, 34050 (street codes). What is the probability that it will get a ticket?

Probability of Black vehicle parking illegally at 34510, 10030, 34050 getting ticket:

[Stage 74:=====> (21 + 1) / 22]

prediction	Count	Total_Cars	probability
2	533813	3575206	0.1493097180973628

NBA Shot Logs

1. For each pair of the players (A, B), we define the fear score of A when facing B as the hit rate, such that B is the closest defender when A is shooting. Based on the fear score, for each player, please find out who is his "most unwanted defender"

player_name	CLOSEST_DEFENDER	hit_rate	total_shots
aaron brooks	Nurkic, Jusuf	0.0	5
aaron gordon	Rivers, Austin	0.0	2
al farouq aminu	Johnson, James	0.0	3
al horford	Diaw, Boris	0.0	4
al jefferson	Hardaway Jr., Tim	0.0	3
alan anderson	Leuer, Jon	0.0	5
alan crabbe	Sefolosha, Thabo	0.0	3
alex len	Knight, Brandon	0.0	4
alexis ajinca	Meeks, Jodie	0.0	3
alonzo gee	Korver, Kyle	0.0	3
amare stoudemire	Deng, Luol	0.0	3
amir johnson	Tucker, PJ	0.0	3
andre drummond	James, LeBron	0.0	5
andre iguodala	Lowry, Kyle	0.0	3
andre miller	Turner, Evan	0.0	2
andre roberston	Ginobili, Manu	0.0	2
andrew bogut	Ibaka, Serge	0.0	6
andrew wiggins	Roberts, Brian	0.0	9
anthony bennett	Ajinca, Alexis	0.0	5
anthony davis	Mbah a Moute, Luc	0.0	5

only showing top 20 rows

2. For each player, we define the comfortable zone of shooting is a matrix of, {SHOT DIST, CLOSE DEF DIST, SHOT CLOCK}. Considering the hit rate, which zone is the best for James Harden, Chris Paul, Stephen Curry, and LeBron James.

player_name	prediction (zone)	Best_zone_res
chris paul	2	0.5531914893617021
james harden	0	0.5604395604395604
lebron james	0	0.6613545816733067
stephen curry	0	0.6350710900473934